



Great Learning

4-5 minutes

Now that we understand what AI Agents are, the different types of AI Agent architectures, and the role and types of AI Agent environments, it's important to note that LLM-based agents are generative by nature and can produce fluent but incorrect or unsupported responses. In autonomous or semi-autonomous systems, such behavior can lead to incorrect decisions, policy violations, or loss of user trust.

We need a mechanism to ensure that

- the agent responses are supported by evidence such as retrieved documents, structured data, or tool outputs
- add a verification layer that checks correctness, consistency, and compliance before responses or actions are finalized

Introduction to Groundedness and Validation

Groundedness in AI agents refers to the practice of anchoring agent outputs to verifiable data, tools, or trusted external context rather than relying solely on the model's internal knowledge.

Validation ensures that the agent's outputs and actions are correct, consistent, and aligned with predefined constraints, policies, and business rules.

Together, groundedness and validation form the foundation for building reliable, trustworthy, and enterprise-ready agentic systems.

Without grounding and validation:

- Agents may hallucinate facts or fabricate explanations
- Outputs may be inconsistent across runs
- Business rules and safety constraints may be violated

Where Groundedness and Validation Are Applied

Groundedness and validation should be implemented at multiple stages of the agent workflow:

- Context Retrieval: Ensuring that responses are generated only from relevant and reliable sources.

